



(a) admission probe — pooled over 4 prompt variants



(b) per prompt variant — mean pooling



*Behavioural pre-commitment probe. For each URL span in the rerank prompt, the label is 1 if the (model, variant) admitted that URL, else 0.*

*Linear probe (logistic regression on frozen hidden states), one fit per (layer, pooling), 80/20 stratified split. Pooled across Llama-3.3-70B and Qwen-2.5-72B. Shaded band in (a): min-max envelope across the 4 prompt variants.*