





Linear probes (logistic regression on frozen hidden states), one fit per (treatment, layer, pooling) tuple, 80/20 stratified split.

Trained on Llama-3.3-70B and Qwen-2.5-72B; results averaged across the four prompt variants {biased, neutral, biased+RAG, neutral+RAG}. Shaded bands in (a): min-max envelope across the 4 variants (invisibly narrow — variant has no measurable effect on the probe).